MIT | Academy of Engineering

EDS ACTIVITY NO. 1

NAME : ANIKET ANIL BHUJAD

DIV : CS6

PRN : 202401080037

ROLL NO. CS6-87

BATCH : C64

- For pandas : -
  - Dataset:

```
[1]: import pandas as pd

data = {
    "Date": pd.date_range(start="2025-01-01", periods=6, freq="D"),
    "Product": ["A", "B", "A", "C", "B", "A"],
    "Units_Sold": [10, 5, 15, 20, 8, 12],
    "Unit_Price": [100, 200, 100, 150, 200, 100]
}
df = pd.DataFrame(data)
df["Revenue"] = df["Units_Sold"] * df["Unit_Price"]

print(df)
```

```
        Date Product  Units_Sold  Unit_Price  Revenue
0 2025-01-01       A          10         100     1000
1 2025-01-02       B           5         200     1000
2 2025-01-03       A          15         100     1500
3 2025-01-04       C          20         150     3000
4 2025-01-05       B           8         200     1600
5 2025-01-06       A          12         100     1200
```

Grains : -

1.Show the first few rows of the dataset

2.Get descriptive statistics for all columns

3.Filter rows where Units_Sold > 10

4.Add a new column 'Tax' as 10% of Revenue

5.Sort rows by Revenue in descending order

6.Find the mean of Units_Sold and Revenue columns

7.Filter rows where Discount is 0

8.Group by Unit_Price and get sum of values

9.Create a Discounted_Revenue column after applying discount

10.Select only the first 3 columns of the dataframe

Outputs in jupyter notebook for all grains: -

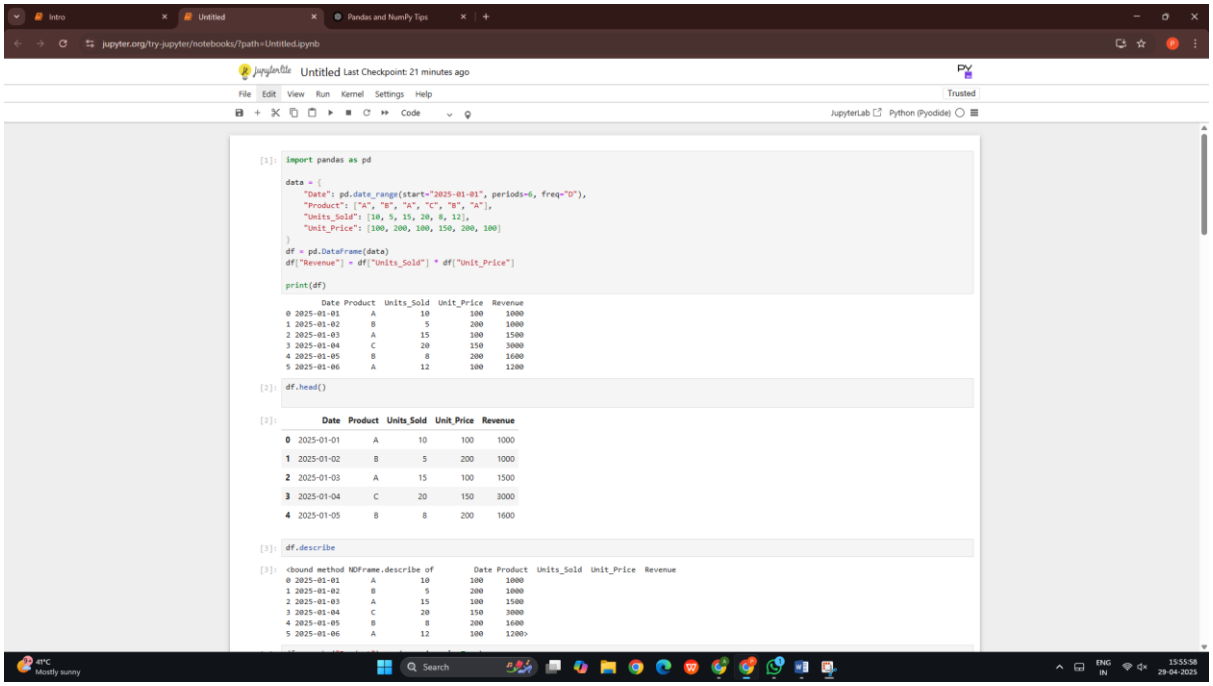

```
import pandas as pd

data = {
    "Date": pd.date_range(start="2025-01-01", periods=6, freq="D"),
    "Product": ["A", "B", "A", "C", "B", "A"],
    "Units_Sold": [10, 5, 15, 20, 8, 12],
    "Unit_Price": [100, 200, 100, 150, 200, 100]
}
df = pd.DataFrame(data)
df["Revenue"] = df["Units_Sold"] * df["Unit_Price"]

print(df)
```

```
        Date Product  Units_Sold  Unit_Price  Revenue
0 2025-01-01       A          10         100     1000
1 2025-01-02       B           5         200     1000
2 2025-01-03       A          15         100     1500
3 2025-01-04       C          20         150     3000
4 2025-01-05       B           8         200     1600
5 2025-01-06       A          12         100     1200
```

```
df.head()
```

	Date	Product	Units_Sold	Unit_Price	Revenue
0	2025-01-01	A	10	100	1000
1	2025-01-02	B	5	200	1000
2	2025-01-03	A	15	100	1500
3	2025-01-04	C	20	150	3000
4	2025-01-05	B	8	200	1600

```
df.describe
```

```
<bound method NDFrame.describe of        Date Product  Units_Sold  Unit_Price  Revenue
0 2025-01-01       A          10         100     1000
1 2025-01-02       B           5         200     1000
2 2025-01-03       A          15         100     1500
3 2025-01-04       C          20         150     3000
4 2025-01-05       B           8         200     1600
5 2025-01-06       A          12         100     1200>
```

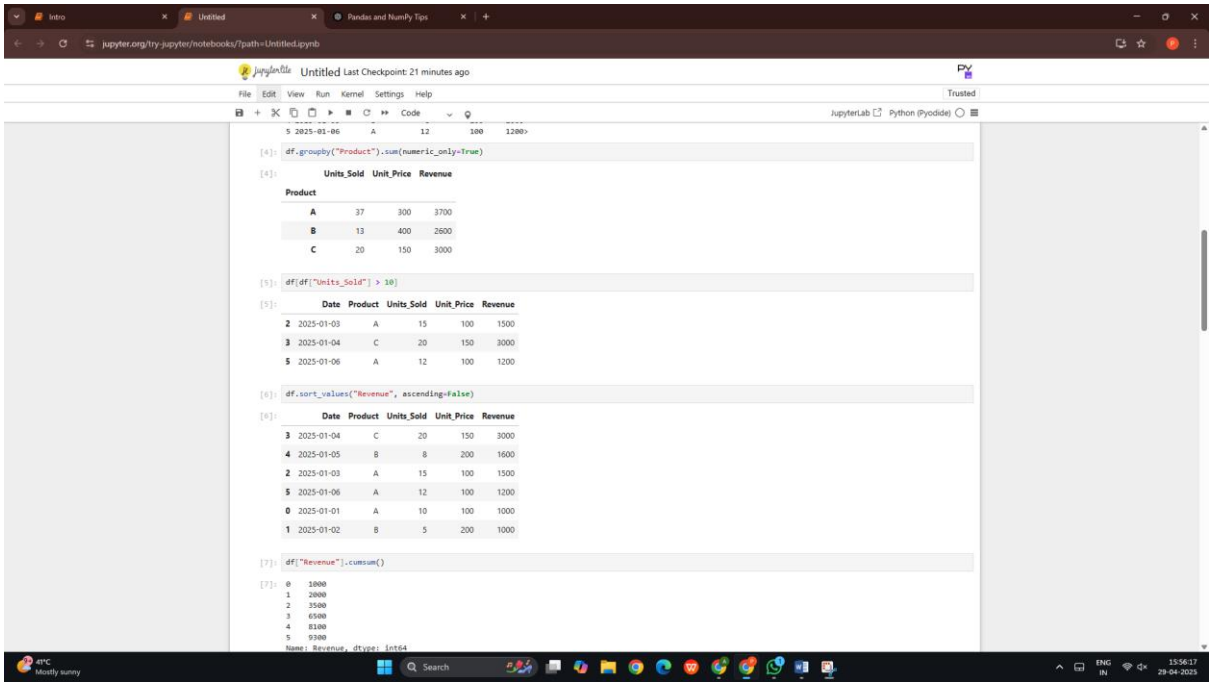

```
df.groupby("Product").sum(numeric_only=True)
```

Product	Units_Sold	Unit_Price	Revenue
A	37	300	3700
B	13	400	2600
C	20	150	3000

```
df[df["Units_Sold"] > 10]
```

	Date	Product	Units_Sold	Unit_Price	Revenue
2	2025-01-03	A	15	100	1500
3	2025-01-04	C	20	150	3000
5	2025-01-06	A	12	100	1200

```
df.sort_values("Revenue", ascending=False)
```

	Date	Product	Units_Sold	Unit_Price	Revenue
3	2025-01-04	C	20	150	3000
4	2025-01-05	B	8	200	1600
2	2025-01-03	A	15	100	1500
5	2025-01-06	A	12	100	1200
0	2025-01-01	A	10	100	1000
1	2025-01-02	B	5	200	1000

```
df["Revenue"].cumsum()
```

```
0    1000
1    2000
2    3500
3    6500
4    8100
5    9300
Name: Revenue, dtype: int64
```

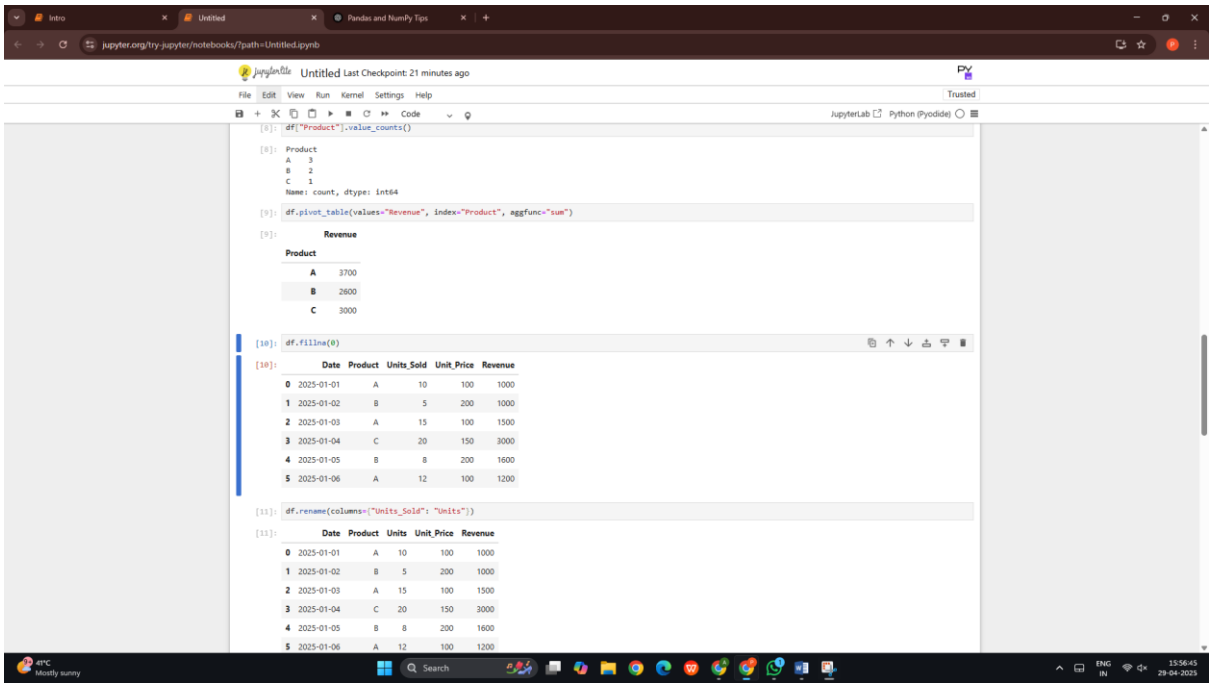
Intro
Untitled
Pandas and NumPy Tips
jupyter.org/try-jupyter/notebooks/?path=Untitled.ipynb
Untitled Last Checkpoint: 21 minutes ago
File Edit View Run Kernel Settings Help
Trusted
Code
JupyterLab
Python (Pyodide)
[8]: df["Product"].value_counts()
[8]: Product
A 3
B 2
C 1
Name: count, dtype: int64
[9]: df.pivot_table(values="Revenue", index="Product", aggfunc="sum")
[9]: Revenue
Product
A 3700
B 2600
C 3000
[10]: df.fillna(0)
[10]: Date Product Units_Sold Unit_Price Revenue
0 2025-01-01 A 10 100 1000
1 2025-01-02 B 5 200 1000
2 2025-01-03 A 15 100 1500
3 2025-01-04 C 20 150 3000
4 2025-01-05 B 8 200 1600
5 2025-01-06 A 12 100 1200
[11]: df.rename(columns={"Units_Sold": "Units"})
[11]: Date Product Units Unit_Price Revenue
0 2025-01-01 A 10 100 1000
1 2025-01-02 B 5 200 1000
2 2025-01-03 A 15 100 1500
3 2025-01-04 C 20 150 3000
4 2025-01-05 B 8 200 1600
5 2025-01-06 A 12 100 1200
41°C
Mostly sunny
Search
ENG
IN
15:56:45
29-04-2025

- For numpy : -
  - Dataset : -

```
[12]: import numpy as np

# 6 rows (sales records), 4 columns (features)
# Columns: Units_Sold, Unit_Price, Revenue, Discount
arr = np.array([
    [10, 100, 1000, 5],
    [5, 200, 1000, 10],
    [15, 100, 1500, 0],
    [20, 150, 3000, 15],
    [8, 200, 1600, 10],
    [12, 100, 1200, 0]
])
```

Grains : -

1. Reshape the array to 3 rows and 8 columns
2. Find the mean of each column
3. Find rows where Units_Sold > 10
4. Add a new column for Tax = 10% of Revenue
5. Transpose the array
6. Find the maximum value in Revenue column
7. Normalize the Units_Sold column
8. Find rows where Discount is 0
9. Find the total Revenue
10. Sort the array by Units_Sold

Output in jupyter notebook for all grains : -